

CONVOLUTIONAL NEURAL NETWORKS: DISCRETE CONVOLUTIONS



Convolution
operations first
published by
D'Alembert in 1754

Discrete convolutions
are matrix operations
that can be used to
apply **filters** to a
matrix or array


The convolutional
neural network
architecture was first
described by
Kunihiko Fukushima
in 1980

pre-defined filters

Discrete convolutions
are matrix operations
that can be used to
apply **filters** to a
matrix or array

Discrete Convolutions $C = A \circledast h$

$$C[m, n] = \sum_{j=-\omega}^{\omega} \sum_{i=-\omega}^{\omega} h[i + \omega, j + \omega] * A[m + i, n + j]$$




Discrete Convolutions $C = A \circledast h$

Filter

Stride


Edges



Discrete Convolutions $C = A \circledast h$

Stride


Edges



Discrete Convolutions $C = A \circledast h$

Stride

Edges





Discrete Convolutions $C = A \circledast h$

Filter
(11x11)

Stride
1

Edges
None




Discrete Convolutions $C = A \circledast h$

Filter
(11x11)

Stride
1

Edges
None





Discrete Convolutions $C = A \circledast h$

Filter
(11x11)



Stride
1

Edges
None



Discrete Convolutions $C = A \circledast h$

Blurring Filter





Discrete Convolutions $C = A \circledast h$

Filter
(3x3)

Stride
1

Edges
None




Discrete Convolutions $C = A \circledast h$

Filter
(3x3)

Stride
1

Edges
None





Discrete Convolutions $C = A \circledast h$

Filter
(3x3)

Stride
1

Edges
None




Discrete Convolutions $C = A \circledast h$

Filter
(3x3)

Stride
1

Edges
None





Discrete Convolutions $C = A \circledast h$

Filter
(3x3)


Stride
1

Edges
None



Discrete Convolutions $C = A \circledast h$

Line Detection / Extraction



Discrete Convolutions $C = A \circledast h$

Photograph







Discrete Convolutions $C = A \circledast h$

Photograph






Discrete Convolutions $C = A \circledast h$

Filter
(101x101)



N-D Discrete Convolutions

$$C[m, n] = \sum_{i=-\omega}^{\omega} h[i + \omega] * A[m + i]$$





N-D Discrete Convolutions

Filter
(1x5)

Stride
1

Edges
None



0.5 0.5 0.5 0.5 0.5


Discrete Convolutions

Discrete convolutions
apply to **array** or
matrix-like data


Discrete convolutions
are matrix operations
that can be used to
apply **filters** to a
matrix or array

These filters can
extract features or
transform the input


CONVOLUTIONAL NEURAL NETWORKS: ARCHITECTURE AND TRAINING




CONVOLUTIONAL NEURAL NETWORKS




INPUTS CONSIDERED "UNORDERED"





BACKPROPAGATION

$$w_1 = w_1 + \eta * \frac{\partial J}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial q_1} \frac{\partial q_1}{\partial w_1}$$




$$w_2 = w_2 + \eta * \frac{\partial J}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial q_2} \frac{\partial q_2}{\partial w_2}$$

Loss function

$$J(w) = \hat{y} - y$$


BACKPROPAGATION

$$\vec{o} = \vec{x} \odot \vec{F}$$






BACKPROPAGATION

$$\vec{o} = \vec{x} \circledast \vec{F}$$



$$o[n] = (x \circledast F)[n] = \sum_{i=-\omega}^{\omega} x[i + n + \omega] * F[i + \omega]$$






BACKPROPAGATION

$$F_i \leftarrow F_i - \eta \frac{\partial J}{\partial F_i}$$

$$\frac{\partial J}{\partial F_i} = \sum_{k=1}^M \frac{\partial J}{\partial o_k} \frac{\partial o_k}{\partial F_i}$$

$$o[n] = (x \circledast F)[n] = \sum_{i=-\omega}^{\omega} x[i + n + \omega] * F[i + \omega]$$




BACKPROPAGATION

$$F_i \leftarrow F_i - \eta \frac{\partial J}{\partial F_i}$$

$$\frac{\partial J}{\partial F_i} = \sum_{k=1}^M \frac{\partial J}{\partial o_k} \frac{\partial o_k}{\partial F_i}$$

$$o[n] = (x \circledast F)[n] = \sum_{i=-\omega}^{\omega} x[i + n + \omega] * F[i + \omega]$$




BACKPROPAGATION

$$\frac{\partial J}{\partial F_1} = \frac{\partial J}{\partial o_0} \frac{\partial o_0}{\partial F_1} + \frac{\partial J}{\partial o_1} \frac{\partial o_1}{\partial F_1}$$

$$F_i \leftarrow F_i - \eta \frac{\partial J}{\partial F_i}$$

$$\frac{\partial J}{\partial F_i} = \sum_{k=1}^M \frac{\partial J}{\partial o_k} \frac{\partial o_k}{\partial F_i}$$

$$o[n] = (x \circledast F)[n] = \sum_{i=-\omega}^{\omega} x[i + n + \omega] * F[i + \omega]$$




BACKPROPAGATION

$$\frac{\partial J}{\partial F_1} = \boxed{\frac{\partial J}{\partial o_0}} \frac{\partial o_0}{\partial F_1} + \boxed{\frac{\partial J}{\partial o_1}} \frac{\partial o_1}{\partial F_1}$$

$$F_i \leftarrow F_i - \eta \frac{\partial J}{\partial F_i}$$

$$\frac{\partial J}{\partial F_i} = \sum_{k=1}^M \frac{\partial J}{\partial o_k} \frac{\partial o_k}{\partial F_i}$$

$$o[n] = (x \circledast F)[n] = \sum_{i=-\omega}^{\omega} x[i + n + \omega] * F[i + \omega]$$



BACKPROPAGATION

$$\frac{\partial J}{\partial F_1} = \frac{\frac{\partial J}{\partial o_0} \frac{\partial o_0}{\partial F_1}}{\frac{\partial J}{\partial o_0} \frac{\partial o_0}{\partial F_1}} + \frac{\frac{\partial J}{\partial o_1} \frac{\partial o_1}{\partial F_1}}{\frac{\partial J}{\partial o_1} \frac{\partial o_1}{\partial F_1}}$$


$$o_0 = F_0 x_0 + F_1 x_1$$

$$o_1 = F_0 x_1 + F_1 x_2$$

$$F_i \leftarrow F_i - \eta \frac{\partial J}{\partial F_i}$$

$$\frac{\partial J}{\partial F_i} = \sum_{k=1}^M \frac{\partial J}{\partial o_k} \frac{\partial o_k}{\partial F_i}$$

$$o[n] = (x \circledast F)[n] = \sum_{i=-\omega}^{\omega} x[i + n + \omega] * F[i + \omega]$$



BACKPROPAGATION

$$\frac{\partial J}{\partial F_1} = \frac{\frac{\partial J}{\partial o_0} \frac{\partial o_0}{\partial F_1}}{\frac{\partial J}{\partial o_0} \frac{\partial o_0}{\partial F_1}} + \frac{\frac{\partial J}{\partial o_1} \frac{\partial o_1}{\partial F_1}}{\frac{\partial J}{\partial o_1} \frac{\partial o_1}{\partial F_1}}$$

$$o_0 = F_0 x_0 + F_1 x_1$$

$$o_1 = F_0 x_1 + F_1 x_2$$


$$\frac{\partial o_0}{\partial F_1} = x_1$$

$$\frac{\partial o_1}{\partial F_1} = x_2$$


$$F_i \leftarrow F_i - \eta \frac{\partial J}{\partial F_i}$$

$$\frac{\partial J}{\partial F_i} = \sum_{k=1}^M \frac{\partial J}{\partial o_k} \frac{\partial o_k}{\partial F_i}$$


$$o[n] = (x \circledast F)[n] = \sum_{i=-\omega}^{\omega} x[i + n + \omega] * F[i + \omega]$$




CONVOLUTIONAL NEURAL NETWORKS



CONVOLUTIONAL NEURAL NETWORKS



CONVOLUTIONAL NEURAL NETWORKS



POOLING


1	1	2	4
5	6	9	3
3	2	4	4
1	2	0	7

max pool with 2x2 filters
and stride 2




6	9
3	7


BACKPROPOGATION




CONVOLUTIONAL NEURAL NETWORKS



CHECK: WHAT HAPPENS WITH TRANSLATION IN INPUT?




CHECK: PROS? CONS?




Application 2: Can we use machine learning to **accurately** classify events in detectors?

Metrics



Detect Lung Cancer

99% Accuracy



TRUE

Proton

Not Proton

PREDICTED

Proton

Not Proton

TRUE
POSITIVE
(TP)

FALSE
NEGATIVE
(FN)

FALSE
POSITIVE
(FP)

TRUE
NEGATIVE
(TN)

TRUE

Proton

Not Proton

PREDICTED			
		Proton	Not Proton
TRUE	Proton	TRUE POSITIVE (TP)	FALSE NEGATIVE (FN)
	Not Proton	FALSE POSITIVE (FP)	TRUE NEGATIVE (TN)

$$\text{accuracy} = \frac{TP + TN}{TP + FN + FP + TN}$$

$$\text{precision} = \frac{TP}{TP + FP}$$

$$\text{recall} = \frac{TP}{TP + FN}$$

$$F1 = \frac{2 \cdot \text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

TRUE

Proton

Not Proton

PREDICTED

Proton

Not Proton

TRUE
POSITIVE
(TP)

TRUE
NEGATIVE
(TN)

PERFECT MODEL

TRUE

Proton

Not Proton

PREDICTED

Proton

Not Proton

TRUE
POSITIVE
(TP)


FALSE
NEGATIVE
(FN)

FALSE
POSITIVE
(FP)

TRUE
NEGATIVE
(TN)

What does this matrix look like for
multi-class classification?


MULTI-CLASS CLASSIFICATION



$$\text{Softmax}(\mathbf{x})_i = \frac{e^{x_i}}{\sum_{j=1}^C e^{x_j}}$$

LOSS FUNCTIONS

Loss function



Mean squared error
(MSE)

$$J(w) = \frac{1}{N} \sum_{i=0}^N (\hat{y}_i - y_i)^2$$

Mean absolute error
(MAE)

$$J(w) = \frac{1}{N} \sum_{i=0}^N |\hat{y}_i - y_i|$$

Cross entropy
(CE)

$$\mathcal{L}(\mathbf{y}, \hat{\mathbf{y}}) = - \sum_{i=1}^N \sum_{j=1}^C y_{ij} \cdot \log(\hat{y}_{ij})$$

Another approach:

Vision Transformers


AN IMAGE IS WORTH 16X16 WORDS: TRANSFORMERS FOR IMAGE RECOGNITION AT SCALE

Alexey Dosovitskiy^{*,†}, Lucas Beyer^{*}, Alexander Kolesnikov^{*}, Dirk Weissenborn^{*},
Xiaohua Zhai^{*}, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer,
Georg Heigold, Sylvain Gelly, Jakob Uszkoreit, Neil Houlsby^{*,†}

^{*}equal technical contribution, [†]equal advising


Google Research, Brain Team

{adosovitskiy, neilhoulshby}@google.com



Vision Transformer

1	0	1	0	0	0	0	0
0	1	1	0	0	0	0	0
0	0	1	0	0	0	0	0
0	0	1	1	0	0	0	0
0	0	1	0	1	0	0	0
1	1	1	1	1	1	1	1
0	0	1	0	0	0	1	0
0	0	1	0	0	0	0	1




Patches

1	0	1	0
0	1	1	0
0	0	1	0
0	0	1	1

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

0	0	1	0
1	1	1	1
0	0	1	0
0	0	1	0

1	0	0	0
1	1	1	1
0	0	1	0
0	0	0	1



Linear "embedding"

0	1	0	1	0	0	1	1	0	0	0	1	0	0	0	1	1
1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
2	0	0	1	0	1	1	1	1	0	0	1	0	0	0	1	0
3	1	0	0	0	1	1	1	1	0	0	1	0	0	0	0	1

Vision Transformer

1	0	1	0	0	0	0	0
0	1	1	0	0	0	0	0
0	0	1	0	0	0	0	0
0	0	1	1	0	0	0	0
0	0	1	0	1	0	0	0
1	1	1	1	1	1	1	1
0	0	1	0	0	0	1	0
0	0	1	0	0	0	0	1

$$X = X_{patch} + P$$

Patches

1	0	1	0
0	1	1	0
0	0	1	0
0	0	1	1

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

0	0	1	0
1	1	1	1
0	0	1	0
0	0	1	0


1	0	0	0
1	1	1	1
0	0	1	0
0	0	0	1

Linear "embedding" with position encoding

0	1	0	1	0	0	1	1	0	0	0	1	0	0	0	1	1
1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
2	0	0	1	0	1	1	1	1	0	0	1	0	0	0	1	0
3	1	0	0	0	1	1	1	1	0	0	1	0	0	0	0	1

Vision Transformer

1	0	1	0	0	0	0	0
0	1	1	0	0	0	0	0
0	0	1	0	0	0	0	0
0	0	1	1	0	0	0	0
0	0	1	0	1	0	0	0
1	1	1	1	1	1	1	1
0	0	1	0	0	0	1	0
0	0	1	0	0	0	0	1




Patches

1	0	1	0
0	1	1	0
0	0	1	0
0	0	1	1

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

0	0	1	0
1	1	1	1
0	0	1	0
0	0	1	0

1	0	0	0
1	1	1	1
0	0	1	0
0	0	0	1



Linear embedding with position

$X =$

0	1	0	1	0	0	1	1	0	0	0	1	0	0	0	1	1
1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
2	0	0	1	0	1	1	1	1	0	0	1	0	0	0	1	0
3	1	0	0	0	1	1	1	1	0	0	1	0	0	0	0	1


Vision Transformer

$X =$

0	1	0	1	0	0	1	1	0	0	0	1	0	0	0	1	1
1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
2	0	0	1	0	1	1	1	1	0	0	1	0	0	0	1	0
3	1	0	0	0	1	1	1	1	0	0	1	0	0	0	0	1


$N \times E$

$W_Q =$




$E \times E$

$W_K =$



$E \times E$

$W_V =$



$E \times E$

Vision Transformer

$X =$


0	1	0	1	0	0	1	1	0	0	0	1	0	0	0	1	1
1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
2	0	0	1	0	1	1	1	1	0	0	1	0	0	0	1	0
3	1	0	0	0	1	1	1	1	0	0	1	0	0	0	0	1

$$Q = XW_Q$$


$$K = XW_K$$

$$V = XW_V$$


$W_Q =$



$W_K =$



$W_V =$



Vision Transformer

$X =$


0	1	0	1	0	0	1	1	0	0	0	1	0	0	0	1	1
1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
2	0	0	1	0	1	1	1	1	0	0	1	0	0	0	1	0
3	1	0	0	0	1	1	1	1	0	0	1	0	0	0	0	1

$$Q = XW_Q \quad (N \times E)$$


$$K = XW_K \quad (N \times E)$$

$$V = XW_V \quad (N \times E)$$


$W_Q =$



$W_K =$




$W_V =$




Vision Transformer


$$Q = XW_Q =$$



$$K = XW_K =$$




$$V = XW_V =$$



$$\text{scores} = QK^T$$

Vision Transformer

$$\text{scores} = QK^T$$




Q



K^T

=



Attention:

$$A = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)$$

Output

$$Z = AV$$

Vision Transformer


$$\text{scores} = QK^T$$

Attention:

$$A = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)$$

Output

$$Z = AV$$




Vision Transformer

Output

$$Z = AV$$




Scaled Dot-Product Attention




Vision Transformer

Multi-Head Attention



Scaled Dot-Product Attention



Attention Is All You Need

Ashish Vaswani*
Google Brain
avaswani@google.com

Noam Shazeer*
Google Brain
noam@google.com

Niki Parmar*
Google Research
nikip@google.com

Jakob Uszkoreit*
Google Research
usz@google.com

Llion Jones*
Google Research
llion@google.com

Aidan N. Gomez* †
University of Toronto
aidan@cs.toronto.edu

Lukasz Kaiser*
Google Brain
lukaszkaizer@google.com

Illia Polosukhin* †
illia.polosukhin@gmail.com

CHOOSING AN ARCHITECTURE

HOW MANY LAYERS?

HOW MANY NODES PER LAYER?

LEARNING RATE

DROPOUT?

WHAT ACTIVATION FUNCTION(S)?


HOW MANY CONVOLUTION LAYERS?

FILTER SIZE?


STRIDE?

POOLING?

PRE-TRAINED MODELS
















PRE-TRAINED MODELS




PRETRAINED MODELS









EXAMPLE WORKFLOW



EXAMPLE WORKFLOW



CONVOLUTIONAL NEURAL NETWORKS



OTHER NETWORK

MECHANISMS:

-POINT CONVOLUTIONS

-SYMMETRIES

